

Chapter 21: Concurrency Control techniques

Database Systems CS203



Outline

- Introduction to Locking
- Binary Locks
- Shared/Exclusive Locks
- Two-phase locking concepts
- Two-phase locking techniques
- Dealing with deadlock and starvation
- Time-stamp based concurrency protocols
- Multiversion concurrency protocol

Introduction

- Concurrency control protocols
 - Set of rules to guarantee serializability
- Two-phase locking protocols
 - Lock data items to prevent concurrent access
- Timestamp
 - Unique identifier for each transaction
- Multiversion concurrency control protocols
 - Use multiple versions of a data item
- Validation or certification of a transaction

21.1 Two-Phase Locking Techniques for Concurrency Control

- Lock
 - Variable associated with a data item describing status for operations that can be applied
 - One lock for each item in the database
- Binary locks
 - Two states (values)
 - Locked (1)
 - Item cannot be accessed
 - Unlocked (0)
 - Item can be accessed when requested

21.1 Two-Phase Locking Techniques for Concurrency Control (cont'd)

- Transaction requests access by issuing a `lock_item(X)` operation

```
lock_item(X):  
  B:  if LOCK(X) = 0                (*item is unlocked*)  
        then LOCK(X) ← 1          (*lock the item*)  
      else  
        begin  
          wait (until LOCK(X) = 0  
                and the lock manager wakes up the transaction);  
          go to B  
        end;  
unlock_item(X):  
  LOCK(X) ← 0;                      (* unlock the item *)  
  if any transactions are waiting  
    then wakeup one of the waiting transactions;
```

Figure 21.1 Lock and unlock operations for binary locks

21.1 Two-Phase Locking Techniques for Concurrency Control (cont'd)

- Lock table specifies items that have locks
- Lock manager subsystem
 - Keeps track of and controls access to locks
 - Rules enforced by lock manager module
- At most one transaction can hold the lock on an item at a given time
- Binary locking too restrictive for database items

21.1 Two-Phase Locking Techniques for Concurrency Control (cont'd)

- Shared/exclusive or read/write locks
 - Read operations on the same item are not conflicting
 - Must have exclusive lock to write
 - Three locking operations
 - `read_lock(X)`
 - `write_lock(X)`
 - `unlock(X)`

Figure 21.2 Locking and unlocking operations for two-mode (read/write, or shared/exclusive) locks

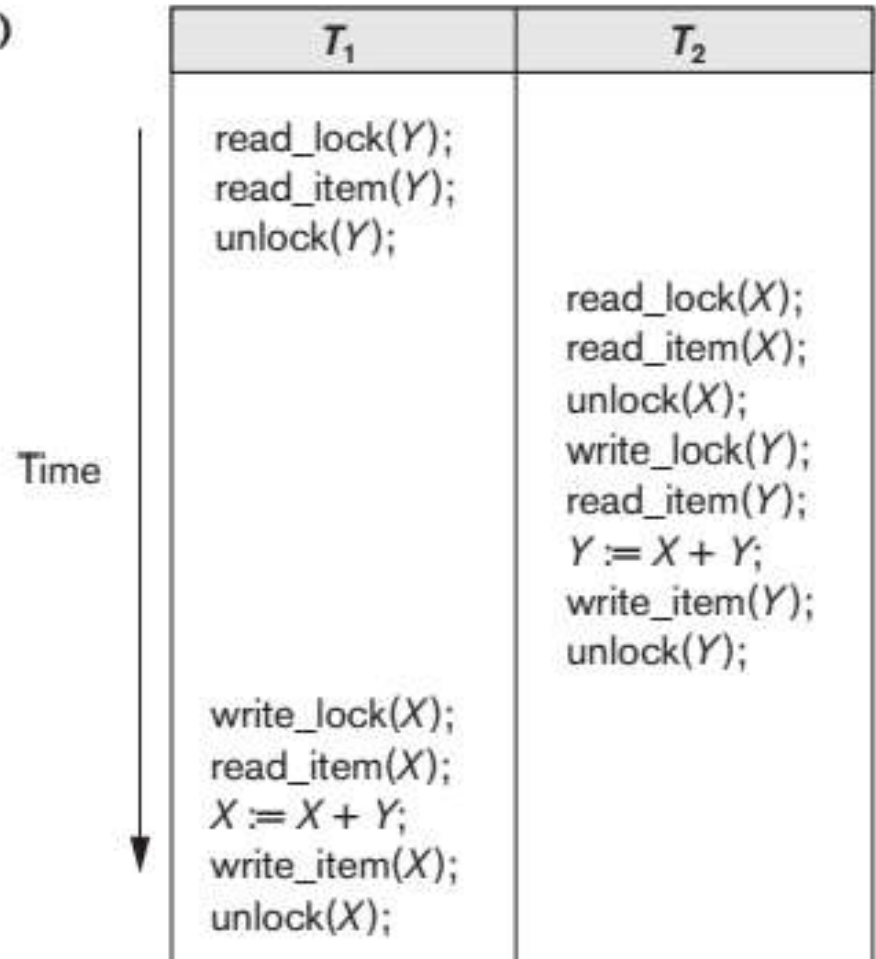
```
read_lock(X):
B:  if LOCK(X) = "unlocked"
    then begin LOCK(X) ← "read-locked";
        no_of_reads(X) ← 1
    end
    else if LOCK(X) = "read-locked"
        then no_of_reads(X) ← no_of_reads(X) + 1
    else begin
        wait (until LOCK(X) = "unlocked"
              and the lock manager wakes up the transaction);
        go to B
    end;

write_lock(X):
B:  if LOCK(X) = "unlocked"
    then LOCK(X) ← "write-locked"
    else begin
        wait (until LOCK(X) = "unlocked"
              and the lock manager wakes up the transaction);
        go to B
    end;

unlock (X):
    if LOCK(X) = "write-locked"
        then begin LOCK(X) ← "unlocked";
            wakeup one of the waiting transactions, if any
        end
    else if LOCK(X) = "read-locked"
        then begin
            no_of_reads(X) ← no_of_reads(X) - 1;
            if no_of_reads(X) = 0
                then begin LOCK(X) = "unlocked";
                    wakeup one of the waiting transactions, if any
                end
            end
        end;
    end;
```


Simple Lock

(c)



21.1 Two-Phase Locking Techniques for Concurrency Control (cont'd)

- Lock conversion
 - Transaction that already holds a lock allowed to convert the lock from one state to another
- Upgrading
 - Issue a read_lock operation then a write_lock operation
- Downgrading
 - Issue a read_lock operation after a write_lock operation

Guaranteeing Serializability by Two-Phase Locking

- Two-phase locking protocol
 - All locking operations precede the first unlock operation in the transaction
 - Phases
 - Expanding (growing) phase
 - New locks can be acquired but none can be released
 - Lock conversion upgrades must be done during this phase
 - Shrinking phase
 - Existing locks can be released but none can be acquired
 - Downgrades must be done during this phase

Guaranteeing Serializability by Two-Phase Locking

- If every transaction in a schedule follows the two-phase locking protocol, schedule guaranteed to be serializable
- Two-phase locking may limit the amount of concurrency that can occur in a schedule
- Some serializable schedules will be prohibited by two-phase locking protocol

Variations of Two-Phase Locking

- Basic 2PL
 - Technique described on previous slides
- Conservative (static) 2PL
 - Requires a transaction to lock all the items it accesses before the transaction begins
 - Predeclare read-set and write-set
 - Deadlock-free protocol
- Strict 2PL
 - Transaction does not release exclusive locks until after it commits or aborts

Variations of Two-Phase Locking (Cont'd)

- Rigorous 2PL
 - Transaction does not release any locks until after it commits or aborts
- Concurrency control subsystem responsible for generating read_lock and write_lock requests
- Locking generally considered to have high overhead

Dealing with Deadlock and Starvation

- Deadlock prevention protocols
 - Every transaction locks all items it needs in advance
 - Ordering all items in the database
 - Transaction that needs several items will lock them in that order
 - Both approaches impractical
- Protocols based on a timestamp
 - Wait-die
 - Wound-wait

Dealing with Deadlock and Starvation (cont'd)

- No waiting algorithm
 - If transaction unable to obtain a lock, immediately aborted and restarted later
- Cautious waiting algorithm
 - Deadlock-free
- Deadlock detection
 - System checks to see if a state of deadlock exists
 - Wait-for graph

Dealing with Deadlock and Starvation (cont'd)

- Victim selection
 - Deciding which transaction to abort in case of deadlock
- Timeouts
 - If system waits longer than a predefined time, it aborts the transaction
- Starvation
 - Occurs if a transaction cannot proceed for an indefinite period of time while other transactions continue normally
 - Solution: first-come-first-served queue

21.2 Concurrency Control Based on Timestamp Ordering

- Timestamp
 - Unique identifier assigned by the DBMS to identify a transaction
 - Assigned in the order submitted
 - Transaction start time
- Concurrency control techniques based on timestamps do not use locks
 - Deadlocks cannot occur

21.2 Concurrency Control Based on Timestamp Ordering (cont'd)

- Generating timestamps
 - Counter incremented each time its value is assigned to a transaction
 - Current date/time value of the system clock
 - Ensure no two timestamps are generated during the same tick of the clock
- General approach
 - Enforce equivalent serial order on the transactions based on their timestamps

21.2 Concurrency Control Based on Timestamp Ordering (cont'd)

- Timestamp ordering (TO)
 - Allows interleaving of transaction operations
 - Must ensure timestamp order is followed for each pair of conflicting operations
- Each database item assigned two timestamp values
 - read_TS(X)
 - write_TS(X)

21.2 Concurrency Control Based on Timestamp Ordering (cont'd)

- Basic TO algorithm
 - If conflicting operations detected, later operation rejected by aborting transaction that issued it
 - Schedules produced guaranteed to be conflict serializable
 - Starvation may occur
- Strict TO algorithm
 - Ensures schedules are both strict and conflict serializable

21.2 Concurrency Control Based on Timestamp Ordering (cont'd)

- Thomas's write rule
 - Modification of basic TO algorithm
 - Does not enforce conflict serializability
 - Rejects fewer write operations by modifying checks for `write_item(X)` operation

21.3 Multiversion Concurrency Control Techniques

- Several versions of an item are kept by a system
- Some read operations that would be rejected in other techniques can be accepted by reading an older version of the item
 - Maintains serializability
- More storage is needed
- Multiversion concurrency control scheme types
 - Based on timestamp ordering
 - Based on two-phase locking
 - Validation and snapshot isolation techniques

21.3 Multiversion Concurrency Control Techniques (cont'd)

- Multiversion technique based on timestamp ordering
 - Two timestamps associated with each version are kept
 - $\text{read_TS}(X_i)$
 - $\text{write_TS}(X_i)$

21.3 Multiversion Concurrency Control Techniques (cont'd)

- Multiversion two-phase locking using certify locks
 - Three locking modes: read, write, and certify

(a)		Read	Write
	Read	Yes	No
	Write	No	No

(b)		Read	Write	Certify
	Read	Yes	Yes	No
	Write	Yes	No	No
	Certify	No	No	No